

# *Fast Random Access Log (FRAL)*

David Maaghul

# Overview

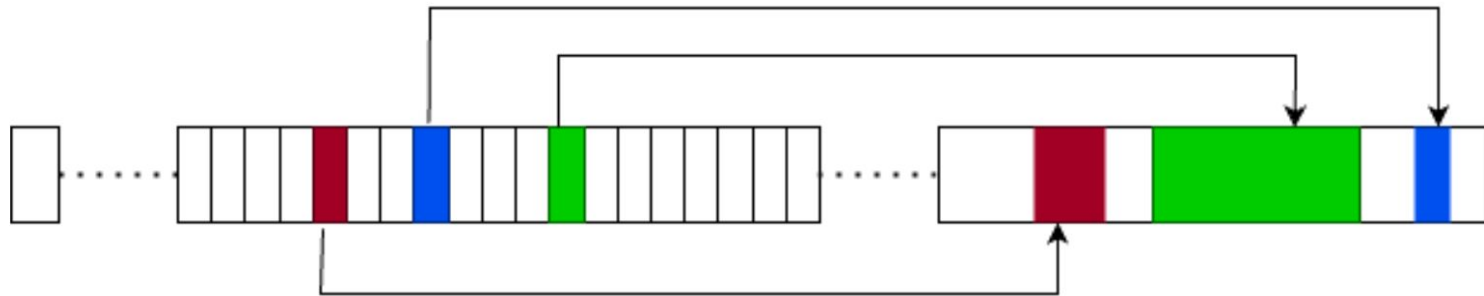
- Goal:
  - Allow multiple processes to read from and write to an array-indexed log as fast as possible
- A Random Access Log:
  - An array-indexed log over shared memory
  - Offers non-blocking writes
  - Reads are random access
- Implementation details:
  - Core engine implemented in C++
    - Python bindings are also provided and tested
- Standard Networking Tools:
  - Implemented and tested tool for syncing logs over network
- Some Use Cases:
  - High Frequency Trading
  - Log Aggregation

# *Core Architecture*

- Three Main Functions:
  - **allocate:** allocates  $n$  contiguous bytes of the memory space (a blob) and returns a reference to the beginning of the allocation. This is similar to malloc in C.
  - **append:** appends a memory allocation (blob) to an indexed log, updates the size of the log and returns the index of the appended blob.
  - **load:** returns the memory allocation of a provided log-entry index

## Core Architecture (Continued)

- Structuring shared memory:
  - Virtual memory = different addressing for different processes
    - Solution: Store offset from start of shared memory allocation in array
  - Array indexed allocations allows for random access of log entries
    - Array entries are random access
  - Atomic variables:
    - Next available address
    - Next available index (estimation)



# *Atomic Allocate*

- Algorithm uses fetch-add primitive
  - Offset of next free memory address is an atomic variable
- Caveats:
  - Atomic offset variable is unbounded and can grow arbitrarily large
  - Process failure after fetch-add step (line 5) results in wasted space.

---

**Algorithm 1** Atomic Allocate

---

```
1:  $s$ : allocation size (bytes)
2:  $S$ : size of memory space (bytes)
3:  $h$ : offset of next free memory address
4: function ALLOCATE( $s$ )
5:    $c = \text{FetchAdd}(h, s)$ 
6:   if  $c + s > S$  then
7:     return null
8:   end if
9:   return  $c$ 
10: end function
```

---

# Atomic Append

- Motivation for approach:
  - What happens if a process fails in the middle of append?
    - Updates atomic index
    - Doesn't index blob
    - Stuck in infinite loop (never break through line 7)
- Want to write an append algorithm that:
  - prioritizes indexing blobs over updating the next-available index variable (atomic)
- Solution:
  - Estimate the next available index

---

## Algorithm 2 Log Reader

---

```
1:  $S$ : max size of log
2: memspace: the shared memory data structure
3: function READ(memspace)
4:    $i \leftarrow 0$ 
5:   for  $i < S$  do
6:      $blob = \text{load}(\text{memspace}, i)$ 
7:     if  $blob$  is not null then
8:       ProcessBlob( $blob$ )
9:        $i \leftarrow i + 1$ 
10:    end if
11:  end for
12: end function
```

---

## Atomic Append (Continued)

- Use compare-and-swap primitive:
  - Only store offset value if corresponding entry is entry
- A process can fail at any point during the append
  - Array entries are never wasted
  - Iterating over log will never yield infinite loop (previous slide)
- In the case of  $n$  writers in parallel:
  - A given writer has a worst case operation count of  $O(n)$
  - The total operation count for

---

### Algorithm 3 Atomic Append

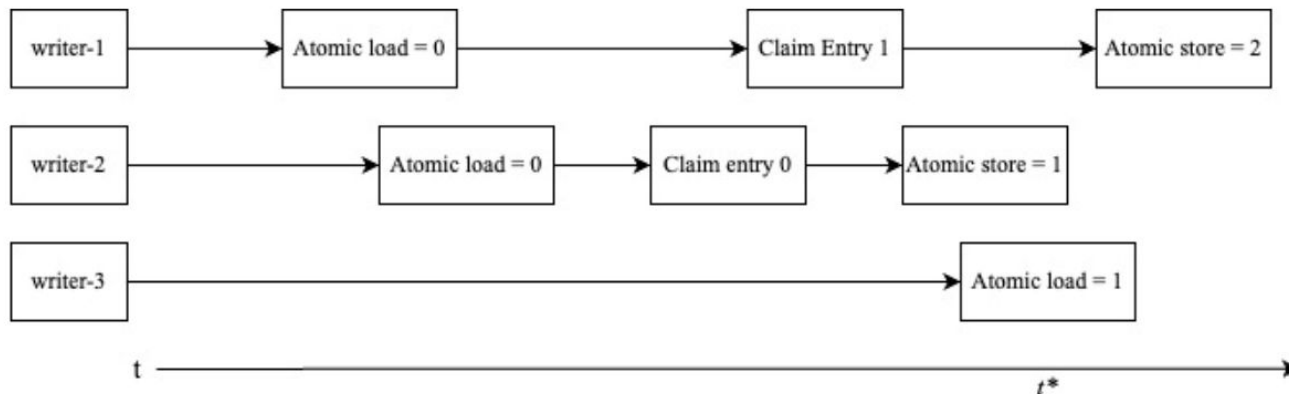
---

```
1: blob: allocation to be appended
2: S: max size of log
3: k: the next available index
4: a: address of start of shared memory space
5: Arr: the log of offsets
6: function APPEND(blob)
7:    $z \leftarrow \text{BlobOffset}(a, \text{blob})$ 
8:    $i \leftarrow \text{AtomicLoad}(k)$ 
9:   for  $i < S$  do
10:    if CompareAndSwap(Arr[i], 0, z) then
11:      AtomicStore(k,  $i + 1$ )
12:      return i
13:    end if
14:     $i \leftarrow i + 1$ 
15:  end for
16:  return -1
17: end function
```

---

## Atomic Append (Continued)

- Problem: can't atomically update two variables at the same time
- The next available index is estimated (relaxed):
  - Example with 3 writers:
    - $k$  = relaxed index of next available entry
    - writer-3 loads  $k = 1$  at time  $t^*$ , but entry 1 has already been claimed by writer-1
    - writer-3 needs two compare-and-swap operations to claim next available index ( $k = 2$ )






# *Performance Testing Results*

- ***Two tests of core architecture performance:***

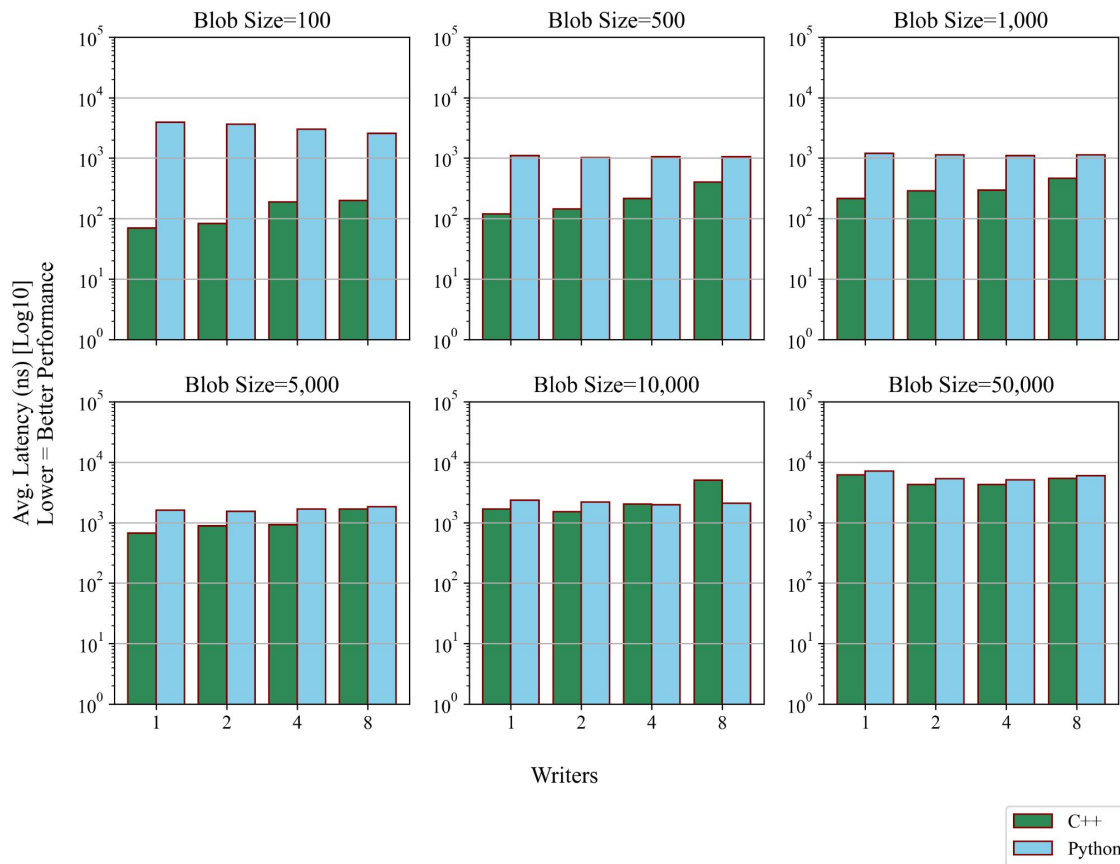
- Look to measure throughput and latency in various ways
- *Write Test:*
  - Measures the speed and efficiency of writing  $N$  GB of data to the log in  $n$ -byte blobs
  - Assesses performance of high volume of write operations by one writer
  - The timer starts before the first write and ends after the last; latency is an average
- *Producer-Consumer Test:*
  - Measures the speed and efficiency of simultaneously writing and reading  $N$  GB of data to the log in  $n$ -bytes blobs with  $k$  writers
  - Assesses performance of concurrent read and write operations from multiple sources
  - The timer starts before the first write and ends after the last read; latency, again, is an average

\*Note that all testing was done with an 8-core Apple M2 machine with 8GB of unified RAM.

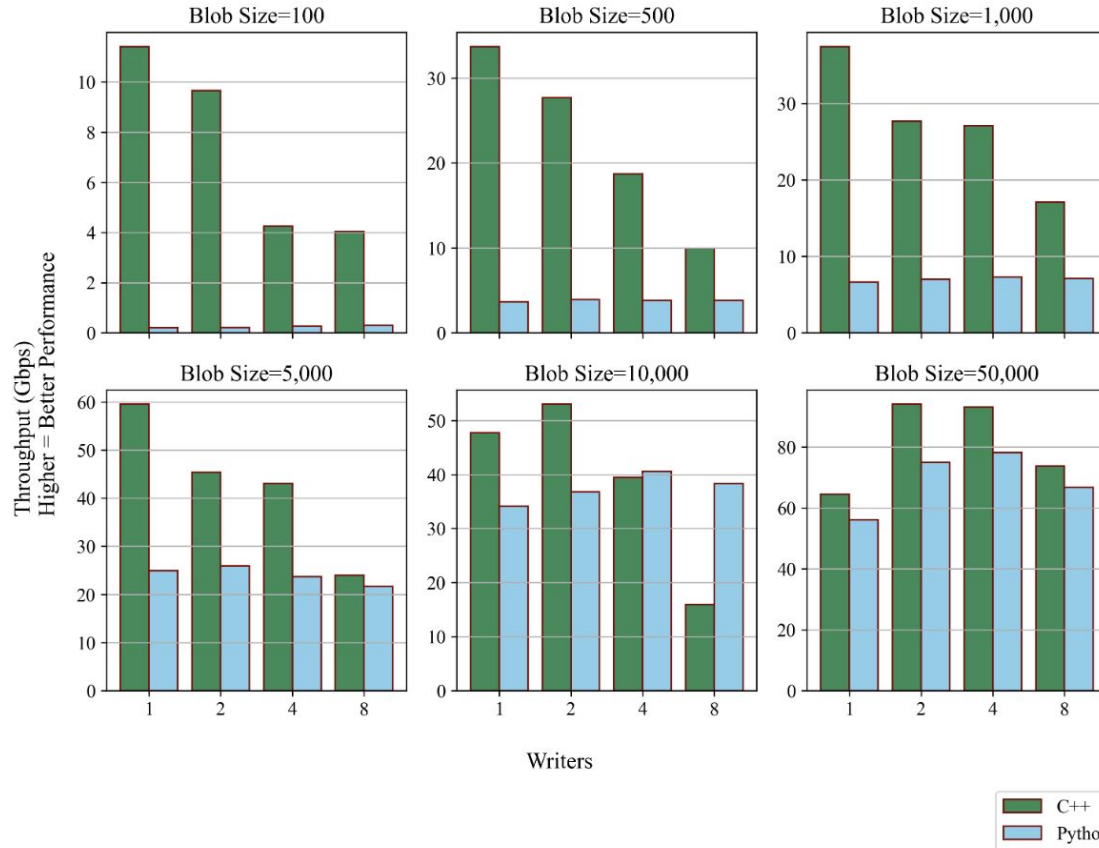
# Write Test Results (Latency + Throughput)



# Producer-Consumer Test Results (Latency)

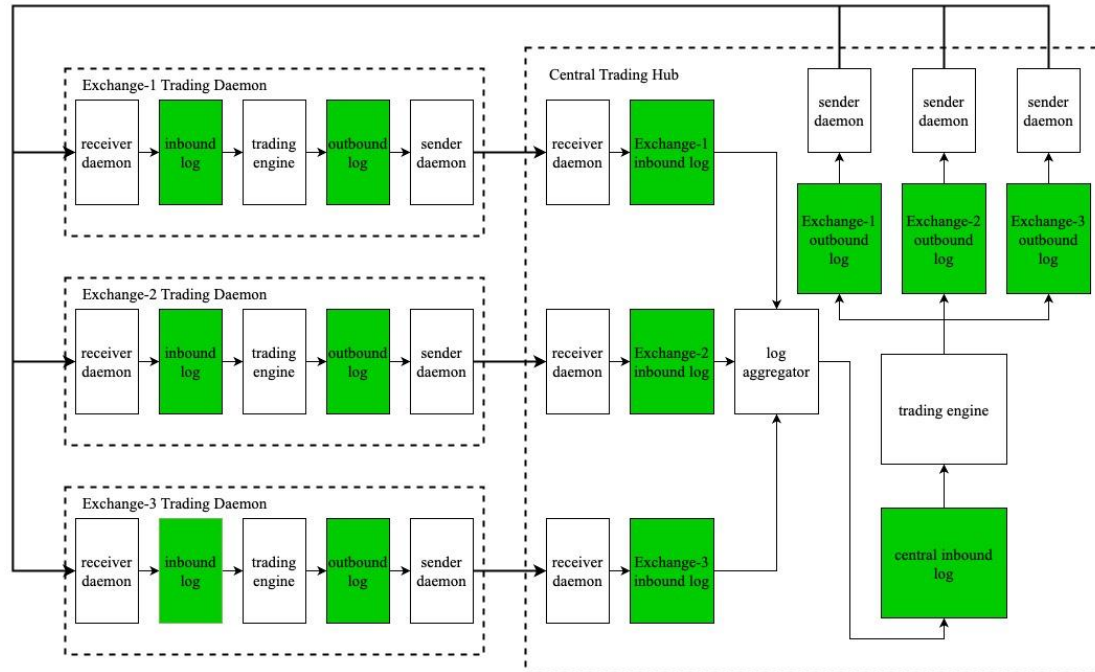


# Producer-Consumer Test Results (Throughput)



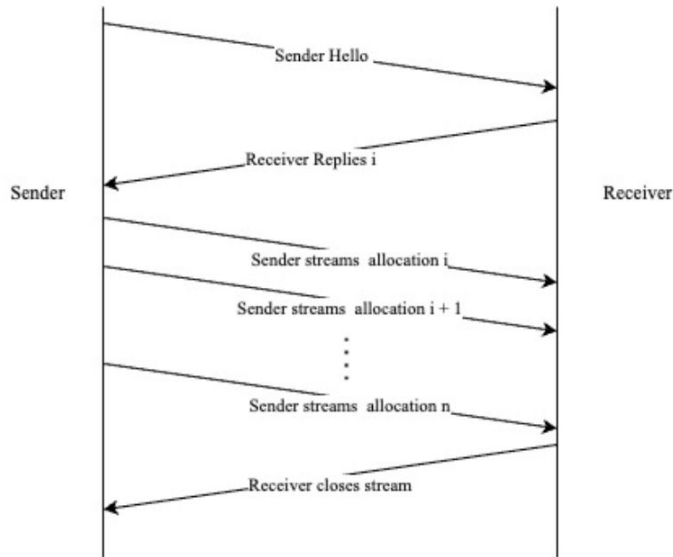
# Syncing Over The Network

- A motivating use case:
  - *High Frequency Trading*: An entity trades on multiple exchanges in different geographic locations, while needing to keep track of net holdings in a central place

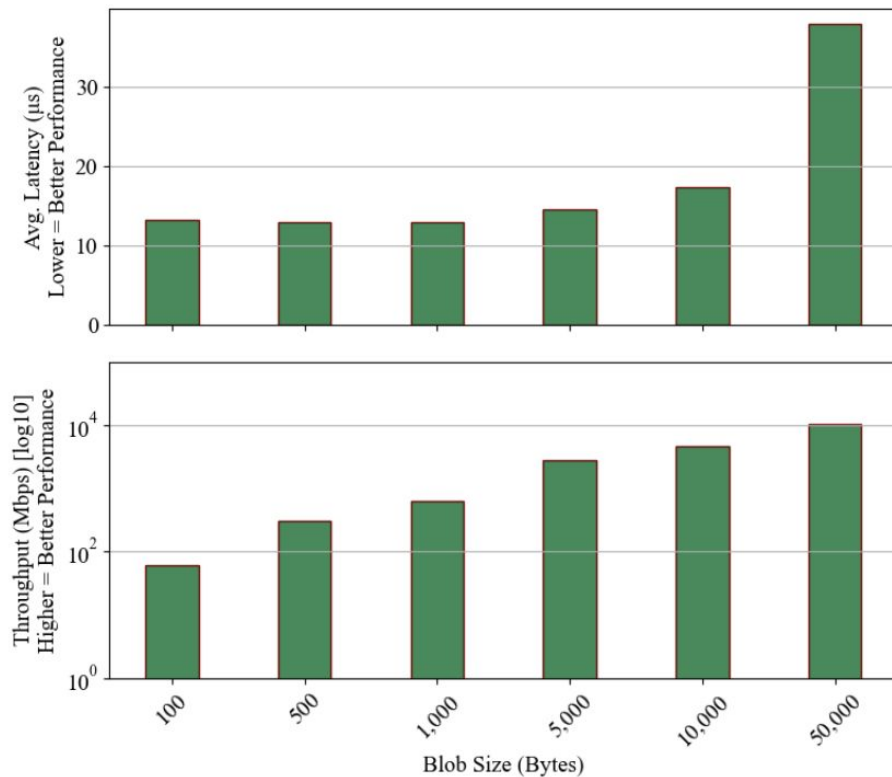


## *Syncing Over The Network (Continued)*

- Intended Requirements:
  - Non-intrusive, i.e., don't need to alter the underlying log structure or add overhead
  - Order of entries is maintained over network
  - Robust to failure for both client and server; mirroring log is identical, through system failures
- Other details:
  - Initial design Implemented over gRPC and Protobuf
  - Future work: non-intrusive network sync tooling from multiple client logs to map to a single client log
- Network Test:
  - Similar to write test, but over the network:
    - Performance degrades (likely to client side buffering by gRPC)
    - Room for improvement



# *Syncing Over The Network (Network Test)*



# Simple Usage (C++)

```
1 //Process 1
2 auto ralA = fral::FRAL("test.bin");
3 const char *TEST_STR = "TEST";
4
5 auto blob = (char *) ralA.allocate(strlen(TEST_STR));
6 strcpy(blob, TEST_STR);
7 printf("%s\n", blob);
8
9 ralA.append(blob);
```

```
1 //Process 2
2 auto ralB = fral::FRAL("test.bin");
3
4 while(true){
5     auto blob2 = (char *) ralB[0]; //load operation
6     if(blob2){
7         break;
8     }
9 }
10
11 printf("%s\n", blob2);
12 blob2[0] = 'D';
13
14 auto blob3 = (char *) ralB[0]; //load operation
15 printf("%s\n", blob3);
```

```
TEST
TEST
DEST
```



# *Simple Usage (Python)*

```
1 # Process 1
2 ral_A = FRAL("test.bin", 1000, 100)
3 test_bytes = "TEST".encode()
4
5 test_blob = ral_A.allocate(len(test_bytes))
6 test_blob[:len(test_bytes)] = test_bytes
7
8 print(bytes(test_blob).decode())
9 ral_A.append(test_blob)
```

```
1 # Process 2
2 ral_B = FRAL("test.bin")
3
4 while True:
5     test_blob2 = ral_B[0] # load operation
6     if test_blob2:
7         break
8
9 print(bytes(test_blob2).decode())
10 test_blob2[0:1] = 'D'.encode()
11
12 test_blob3 = ral_B[0] # load operation
13 print(bytes(test_blob3).decode())
```

```
TEST
TEST
DEST
```

# *Language-Agnostic Demo*

- Language-Agnostic Framework
  - Shared memory opens door for language-agnostic data structures
  - Good for systems with microservices that vary in performance constraints
    - Use C++ for computationally-intensive tasks
    - Capitalize on rapid prototyping in development with Python
    - Example: HFT back-office tasks in Python, trading in C++
  - Demo (video):
    - Read integer from one language, increment by one and write
    - Python initiates process by writing and sending (appending 1)
    - Standard output for writing 5 entries:

```
Python sending 1 to C++!  
C++ received 1 from Python, sending 2!  
Python received 2 from C++, sending 3!  
C++ received 3 from Python, sending 4!  
Python received 4 from C++, sending 5!  
C++ received 5 from Python, done!
```